

Mudcard

- **should we generally stay away from imputation and use onehotencoding instead?**
 - for categorical features, yes
 - If there are missing values in ordinal features, use the ordinal encoder
 - If there are missing values in continuous features, use one of the techniques from today's lecture
- **If the imputation uncertainty is high, what should we do?**
 - Don't impute :)
- **for multivariate imputation, what if all columns are all missing, which should be use to predict the missing values**
 - If literally all values are missing in a column, drop it
- **Not related to the lecture, but should we use IterativeImputer for ps8 2b?**
 - I don't think there are missing values in the diabetes dataset so why would you use it?
- **So if I understand it correctly, the original target variable , 'SalePrice' now becomes one of the feature columns to train the model for predicting the missing value of a feature that now becomes the target variable?**
 - I don't think so
 - If I remember correctly, the columns in the feature matrix only are used in the iterative imputer
 - Please check the publication to be certain
- **I'm wondering why SimpleImputer exists if it is not a good practice, or if there are any good uses to it.**
 - We use the simple imputer to replace NaNs with a string so it is still a useful method and good to have it in sklearn
 - Many people want to do mean or median imputation without realizing that it's not good practice
 - Hopefully I convinced you that imputation is not great and now you know better
- **what does it mean that an imputed data point is "inaccurate"? is there ground truth if we're just sampling randomly from some distribution?**
 - I'm not saying imputed values are inaccurate
 - My point is that we simply don't have a way to measure how accurate the imputed values are because we don't know the ground truth for those missing values
- **Let's say we are observing the same units over time, but some units are only in a few of the later/earlier periods and are missing in others. Is there a point at which it would make more sense to drop them?**

- It kinda depends on why the values are missing, how many values are missing, what the goal of the project is, etc.
- I'd need to know more about the specifics

Missing data, part 2

By the end of this lecture, you will be able to

- apply XGBoost to a dataset with missing values
- apply the reduced-features model (also called the pattern submodel approach)
- decide which approach is best for your dataset

We continue working with the house price data set

- regression problem
- categorical, ordinal, continuous features
- missing data in all feature types

```
In [1]: # read the data
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

# Let's load the data
df = pd.read_csv('../data/train.csv')
# drop the ID
df.drop(columns=['Id'], inplace=True)

# the target variable
y = df['SalePrice']
df.drop(columns=['SalePrice'], inplace=True)
# the unprocessed feature matrix
X = df.values
print(X.shape)
# the feature names
ftrs = df.columns
```

(1460, 79)

```
In [2]: perc_missing_per_ftr = df.isnull().sum(axis=0)/df.shape[0]
print('fraction of missing values in features:')
print(perc_missing_per_ftr[perc_missing_per_ftr > 0])
print('data types of the features with missing values:')
print(df[perc_missing_per_ftr[perc_missing_per_ftr > 0].index].dtypes)
frac_missing = sum(df.isnull().sum(axis=1)!=0)/df.shape[0]
print('fraction of points with missing values:', frac_missing)
```

fraction of missing values in features:

```
LotFrontage    0.177397
Alley          0.937671
MasVnrType     0.597260
MasVnrArea     0.005479
BsmtQual       0.025342
BsmtCond       0.025342
BsmtExposure   0.026027
BsmtFinType1   0.025342
BsmtFinType2   0.026027
Electrical     0.000685
FireplaceQu    0.472603
GarageType     0.055479
GarageYrBlt    0.055479
GarageFinish   0.055479
GarageQual     0.055479
GarageCond     0.055479
PoolQC         0.995205
Fence          0.807534
MiscFeature    0.963014
```

dtype: float64

data types of the features with missing values:

```
LotFrontage    float64
Alley          object
MasVnrType     object
MasVnrArea     float64
BsmtQual       object
BsmtCond       object
BsmtExposure   object
BsmtFinType1   object
BsmtFinType2   object
Electrical     object
FireplaceQu    object
GarageType     object
GarageYrBlt    float64
GarageFinish   object
GarageQual     object
GarageCond     object
PoolQC         object
Fence          object
MiscFeature    object
```

dtype: object

fraction of points with missing values: 1.0

```
In [3]: # let's split to train, CV, and test
X_other, X_test, y_other, y_test = train_test_split(df, y, test_size=0.2, ra
X_train, X_CV, y_train, y_CV = train_test_split(X_other, y_other, test_size=

print(X_train.shape)
print(X_CV.shape)
print(X_test.shape)
```

```
(876, 79)
```

```
(292, 79)
```

```
(292, 79)
```

```
In [4]: # collect the various features
cat_ftrs = ['MSZoning', 'Street', 'Alley', 'LandContour', 'LotConfig', 'Neighborhood',
            'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd',
            'Heating', 'CentralAir', 'Electrical', 'GarageType', 'PavedDrive', 'MiscFeature']
ordinal_ftrs = ['LotShape', 'Utilities', 'LandSlope', 'ExterQual', 'ExterCond', 'Foundation',
               'BsmtFinType1', 'BsmtFinType2', 'HeatingQC', 'KitchenQual', 'FunctionalUse',
               'GarageQual', 'GarageCond', 'PoolQC', 'Fence']
ordinal_cats = [['Reg', 'IR1', 'IR2', 'IR3'], ['AllPub', 'NoSewr', 'NoSeWa', 'ELO'],
               ['Po', 'Fa', 'TA', 'Gd', 'Ex'], ['Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'No', 'Mn', 'Av', 'Gd'],
               ['NA', 'Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'No', 'Mn', 'Av', 'Gd'],
               ['NA', 'Unf', 'LwQ', 'Rec', 'BLQ', 'ALQ', 'GLQ'], ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
               ['Sal', 'Sev', 'Maj2', 'Maj1', 'Mod', 'Min2', 'Min1', 'Typ'], ['NA', 'No', 'Mn', 'Av', 'Gd'],
               ['NA', 'Unf', 'RFn', 'Fin'], ['NA', 'Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'No', 'Mn', 'Av', 'Gd'],
               ['NA', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'MnWw', 'GdWo', 'MnPrv', 'GdPrv']]
num_ftrs = ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
            'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
            'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
            'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars',
            'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea']
```

```
In [5]: # preprocess with pipeline and columntransformer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
import warnings
warnings.filterwarnings("ignore")

# one-hot encoder
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(sparse_output=False, handle_unknown='ignore'))])

# ordinal encoder
ordinal_transformer = Pipeline(steps=[
    ('imputer2', SimpleImputer(strategy='constant', fill_value='NA')),
    ('ordinal', OrdinalEncoder(categories = ordinal_cats))])

# standard scaler
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())])

# collect the transformers
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_ftrs),
        ('cat', categorical_transformer, cat_ftrs),
        ('ord', ordinal_transformer, ordinal_ftrs)])
```

```
In [6]: # fit_transform the training set
X_prep = preprocessor.fit_transform(X_train)
# collect feature names
feature_names = preprocessor.get_feature_names_out()
```

```

df_train = pd.DataFrame(data=X_prep, columns=feature_names)
print(df_train.shape)

# transform the CV
df_cv = preprocessor.transform(X_cv)
df_cv = pd.DataFrame(data=df_cv, columns = feature_names)
print(df_cv.shape)

# transform the test
df_test = preprocessor.transform(X_test)
df_test = pd.DataFrame(data=df_test, columns = feature_names)
print(df_test.shape)

```

```

(876, 220)
(292, 220)
(292, 220)

```

```

In [7]: print('data dimensions:', df_train.shape)
perc_missing_per_ftr = df_train.isnull().sum(axis=0)/df_train.shape[0]
print('fraction of missing values in features:')
print(perc_missing_per_ftr[perc_missing_per_ftr > 0])
frac_missing = sum(df_train.isnull().sum(axis=1)!=0)/df_train.shape[0]
print('fraction of points with missing values:', frac_missing)

```

```

data dimensions: (876, 220)
fraction of missing values in features:
num__LotFrontage    0.173516
num__MasVnrArea     0.004566
num__GarageYrBlt    0.050228
dtype: float64
fraction of points with missing values: 0.2237442922374429

```

Missing data, part 2

By the end of this lecture, you will be able to

- **apply XGBoost to a dataset with missing values**
- apply the reduced-features model (also called the pattern submodel approach)
- decide which approach is best for your dataset

XGBoost

- eXtreme Gradient Boosting - a popular tree-based method
- [blog post](#) and [paper](#)
- more advanced than random forest
 - it has l1 and l2 regularization while random forest does not
 - trees are not independent
 - the next tree is built to improve the previous tree
 - less trees are necessary to achieve same accuracy

- but XGBoost trees can overfit
- handles missing values well

XGBoost and missing values

- some sklearn models raise an error if the feature matrix (X) contains nans - but see [here](#)
- XGBoost doesn't!
- If a feature with missing values is split:
 - XGBoost tries to put the points with missing values to the left and right
 - calculates a metric called gain for both options
 - puts the points with missing values to the side with the higher gain
- if missingness correlates with the target variable, XGBoost extracts this info!

```
In [8]: import xgboost
from sklearn.model_selection import ParameterGrid
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

param_grid = {"learning_rate": [0.03],
              "n_estimators": [10000],
              "seed": [0],
              #"reg_alpha": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
              #"reg_lambda": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
              "missing": [np.nan],
              #"max_depth": [1,3,10,30,100],
              "colsample_bytree": [0.9],
              "subsample": [0.66]}

XGB = xgboost.XGBRegressor()
XGB.set_params(**ParameterGrid(param_grid)[0],early_stopping_rounds=50) # ON
XGB.fit(df_train,y_train,eval_set=[(df_CV, y_CV)], verbose=False)
y_CV_pred = XGB.predict(df_CV)
print('the CV RMSE:',np.sqrt(mean_squared_error(y_CV,y_CV_pred)))
y_test_pred = XGB.predict(df_test)
print('the test RMSE:',np.sqrt(mean_squared_error(y_test,y_test_pred)))
print('the test R2:',r2_score(y_test,y_test_pred))
```

```
the CV RMSE: 24584.178326720623
the test RMSE: 33324.32624975335
the test R2: 0.8391927480697632
```

XGB notes

- do not tune the number of trees, set it to a very large value like 10000
- use early stopping instead! it will automatically determine the best number of trees based on a validation set.
- there are a large number of hyperparameters in XGB
 - tune maybe `max_depth`, `reg_alpha`, and `reg_lambda`

- the default values of some hyperparameters are not optimal
 - set `colsample_bytree` and `subsample` to values a bit smaller than 1 to avoid overfitting

Quiz 1

Missing data, part 2

By the end of this lecture, you will be able to

- apply XGBoost to a dataset with missing values
- apply the reduced-features model (also called the pattern submodel approach)**
- decide which approach is best for your dataset

Reduced-features model (or pattern submodel approach)

- first described in 2007 in a [JMLR article](#) as the reduced features model
- in 2018, "rediscovered" as the pattern submodel approach in [Biostatistics](#)

My test set:

index	feature 1	feature 2	feature 3	target var
0	NA	45	NA	0
1	NA	NA	8	1
2	12	6	34	0
3	1	89	NA	0
4	0	NA	47	1
5	687	24	67	1
6	NA	23	NA	1

To predict points 0 and 6, I will use train and CV points that are complete in feature 2.

To predict point 1, I will use train and CV points that are complete in feature 3.

To predict point 2 and 5, I will use train and CV points that are complete in features 1-3.

Etc. We will train as many models as the number of patterns in test/deployment.

How to determine the patterns?

```
In [9]: mask = df_test[['num__LotFrontage', 'num__MasVnrArea', 'num__GarageYrBlt']].isna()

unique_rows, counts = np.unique(mask, axis=0, return_counts=True)
print(unique_rows.shape) # 6 patterns, we will train 6 models
for i in range(len(counts)):
    print(unique_rows[i], counts[i])
```

(6, 3)

```
[False False False] 223
[False False  True] 21
[False  True False] 1
[ True False False] 44
[ True False  True] 2
[ True  True False] 1
```

```
In [10]: import xgboost
from sklearn.model_selection import ParameterGrid
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

def xgb_model(X_train, Y_train, X_CV, y_CV, X_test, y_test, verbose=1):

    # make into row vectors to avoid an sklearn/xgb warning
    Y_train = np.reshape(np.array(Y_train), (1, -1)).ravel()
    y_CV = np.reshape(np.array(y_CV), (1, -1)).ravel()
    y_test = np.reshape(np.array(y_test), (1, -1)).ravel()

    XGB = xgboost.XGBRegressor(n_jobs=1)

    # find the best parameter set
    param_grid = {"learning_rate": [0.03],
                  "n_estimators": [10000],
                  "seed": [0],
                  #"reg_alpha": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
                  #"reg_lambda": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
                  "missing": [np.nan],
                  #"max_depth": [1,3,10,30,100,],
                  "colsample_bytree": [0.9],
                  "subsample": [0.66]}

    pg = ParameterGrid(param_grid)

    scores = np.zeros(len(pg))

    for i in range(len(pg)):
        if verbose >= 5:
            print("Param set " + str(i + 1) + " / " + str(len(pg)))
        params = pg[i]
        XGB.set_params(**params, early_stopping_rounds=50)
        eval_set = [(X_CV, y_CV)]
        XGB.fit(X_train, Y_train,
                eval_set=eval_set, verbose=False)# with early stopping
        y_CV_pred = XGB.predict(X_CV)
        scores[i] = mean_squared_error(y_CV, y_CV_pred)

    best_params = np.array(pg)[scores == np.max(scores)]
```

```

if verbose >= 4:
    print('Test set max score and best parameters are:')
    print(np.max(scores))
    print(best_params)

# test the model on the test set with best parameter set
XGB.set_params(**best_params[0],early_stopping_rounds=50)
XGB.fit(X_train,Y_train,
        eval_set=eval_set, verbose=False)
y_test_pred = XGB.predict(X_test)

if verbose >= 1:
    print ('The MSE is:',mean_squared_error(y_test,y_test_pred))
if verbose >= 2:
    print ('The predictions are:')
    print (y_test_pred)
if verbose >= 3:
    print("Feature importances:")
    print(XGB.feature_importances_)

return (mean_squared_error(y_test,y_test_pred), y_test_pred, XGB.feature

# Function: Reduced-feature XGB model
# all the inputs need to be pandas DataFrames
def reduced_feature_xgb(X_train, Y_train, X_CV, y_CV, X_test, y_test):

    # find all unique patterns of missing value in test set
    mask = X_test.isnull()
    unique_rows = np.array(np.unique(mask, axis=0))
    all_y_test_pred = pd.DataFrame()

    print('there are', len(unique_rows), 'unique missing value patterns.')

    # divide test sets into subgroups according to the unique patterns
    for i in range(len(unique_rows)):
        print ('working on unique pattern', i)
        ## generate X_test subset that matches the unique pattern i
        sub_X_test = pd.DataFrame()
        sub_y_test = pd.Series(dtype=float)
        for j in range(len(mask)): # check each row in mask
            row_mask = np.array(mask.iloc[j])
            if np.array_equal(row_mask, unique_rows[i]): # if the pattern ma

                sub_X_test = pd.concat([sub_X_test,X_test.iloc[[j]]])# append
                sub_y_test = pd.concat([sub_y_test, y_test.iloc[[j]]])# append

        sub_X_test = sub_X_test[X_test.columns[~unique_rows[i]]]

        ## choose the corresponding reduced features for subgroups
        sub_X_train = pd.DataFrame()
        sub_Y_train = pd.DataFrame()
        sub_X_CV = pd.DataFrame()
        sub_y_CV = pd.DataFrame()
        # 1. remove the feature columns that have nans in the corresponding
        sub_X_train = X_train[X_train.columns[~unique_rows[i]]]
        sub_X_CV = X_CV[X_CV.columns[~unique_rows[i]]]

```

```

# 2.remove the rows in the sub_X_train and sub_X_CV that have any na
sub_X_train = sub_X_train.dropna()
sub_X_CV = sub_X_CV.dropna()
# 3.remove the sub_Y_train and sub_y_CV accordingly
sub_Y_train = Y_train.iloc[sub_X_train.index]
sub_y_CV = y_CV.iloc[sub_X_CV.index]

# run XGB
sub_y_test_pred = xgb_model(sub_X_train, sub_Y_train, sub_X_CV,
                             sub_y_CV, sub_X_test, sub_y_test, ver
sub_y_test_pred = pd.DataFrame(sub_y_test_pred[1], columns=['sub_y_te
                             index=sub_y_test.index)
print('    RMSE:', np.sqrt(mean_squared_error(sub_y_test, sub_y_test_pr
# collect the test predictions
all_y_test_pred = pd.concat([all_y_test_pred, sub_y_test_pred])

# rank the final y_test_pred according to original y_test index
all_y_test_pred = all_y_test_pred.sort_index()
y_test = y_test.sort_index()

# get global RMSE
total_RMSE = np.sqrt(mean_squared_error(y_test, all_y_test_pred))
total_R2 = r2_score(y_test, all_y_test_pred)
return total_RMSE, total_R2

```

```

In [11]: RMSE, R2 = reduced_feature_xgb(df_train, y_train, df_CV, y_CV, df_test, y_te
print('final RMSE:', RMSE)
print('final R2:', R2)

```

there are 6 unique missing value patterns.

working on unique pattern 0

RMSE: 36632.93897071807

working on unique pattern 1

RMSE: 12720.122430512318

working on unique pattern 2

RMSE: 6195.65625

working on unique pattern 3

RMSE: 20333.757172762893

working on unique pattern 4

RMSE: 22206.201991008696

working on unique pattern 5

RMSE: 49318.03125

final RMSE: 33326.26362943341

final R2: 0.8391740548538331

Quiz 2

Missing data, part 2

By the end of this lecture, you will be able to

- apply XGBoost to a dataset with missing values

- apply the reduced-features model (also called the pattern submodel approach)
- **decide which approach is best for your dataset**

Which approach is best for my data?

- **XGB**: run n XGB models with n different seeds
- **imputation**: prepare n different imputations and run n models on them
- **reduced-features**: run n reduced-features model with n different seeds
- rank the three methods based on how significantly different the corresponding mean scores are

Mudcard

In []: